

A JIT compiler for CPython

Brandt Bucher (October 10th, 2023)

Background

Background

- CPython 3.11:
 - Specializing adaptive interpreter profiles programs and optimizes them on-the-fly
- CPython 3.12:
 - Interpreter is generated from a DSL, allowing analysis and modification at build time
- CPython 3.13:
 - Internal pipeline for detecting, optimizing, and executing hot code paths

Background

```
def fibonacci(n):  
    a, b = 0, 1  
    for _ in range(n):  
        a, b = b, a + b  
    return a
```

Background

```
for _ in range(n):  
    a, b = b, a + b
```

Background

Bytecode

```
for _ in range(n):  
    a, b = b, a + b
```

Background

Bytecode

```
for _ in range(n):  
    a, b = b, a + b
```

```
FOR_ITER  
STORE_FAST  
LOAD_FAST_LOAD_FAST  
LOAD_FAST  
BINARY_OP  
STORE_FAST_STORE_FAST  
JUMP_BACKWARD
```

Background

Bytecode

FOR_ITER

STORE_FAST

LOAD_FAST_LOAD_FAST

LOAD_FAST

BINARY_OP

STORE_FAST_STORE_FAST

JUMP_BACKWARD

Background

CPython 3.11: Specialized Bytecode

FOR_ITER

STORE_FAST

LOAD_FAST_LOAD_FAST

LOAD_FAST

BINARY_OP

STORE_FAST_STORE_FAST

JUMP_BACKWARD

Background

CPython 3.11: Specialized Bytecode

FOR_ITER

STORE_FAST

LOAD_FAST_LOAD_FAST

LOAD_FAST

BINARY_OP

STORE_FAST_STORE_FAST

JUMP_BACKWARD

Background

CPython 3.11: Specialized Bytecode

FOR_ITER_RANGE

STORE_FAST

LOAD_FAST_LOAD_FAST

LOAD_FAST

BINARY_OP_ADD_INT

STORE_FAST_STORE_FAST

JUMP_BACKWARD

Background

CPython 3.11: Specialized Bytecode

FOR_ITER_RANGE

STORE_FAST

LOAD_FAST_LOAD_FAST

LOAD_FAST

BINARY_OP_ADD_INT

STORE_FAST_STORE_FAST

JUMP_BACKWARD

Background

CPython 3.13: Micro-Op Traces

FOR_ITER_RANGE

STORE_FAST

LOAD_FAST_LOAD_FAST

LOAD_FAST

BINARY_OP_ADD_INT

STORE_FAST_STORE_FAST

JUMP_BACKWARD

Background

CPython 3.13: Micro-Op Traces (`-X uops`)

FOR_ITER_RANGE

STORE_FAST

LOAD_FAST_LOAD_FAST

LOAD_FAST

BINARY_OP_ADD_INT

STORE_FAST_STORE_FAST

JUMP_BACKWARD

Background

CPython 3.13: Micro-Op Traces (`-X uops`)

FOR_ITER_RANGE

LOAD_FAST

BINARY_OP_ADD_INT

STORE_FAST

STORE_FAST_STORE_FAST

LOAD_FAST_LOAD_FAST

JUMP_BACKWARD

Background

CPython 3.13: Micro-Op Traces (`-X uops`)

FOR_ITER_RANGE

LOAD_FAST

BINARY_OP_ADD_INT

STORE_FAST

STORE_FAST_STORE_FAST

LOAD_FAST_LOAD_FAST

JUMP_BACKWARD

Background

CPython 3.13: Micro-Op Traces (`-X uops`)

```
_SET_IP  
_ITER_CHECK_RANGE  
_IS_ITER_EXHAUSTED_RANGE  
_POP_JUMP_IF_TRUE  
_ITER_NEXT_RANGE
```

```
_SET_IP  
STORE_FAST
```

```
_SET_IP  
LOAD_FAST  
LOAD_FAST
```

```
_SET_IP  
LOAD_FAST  
  
_SET_IP  
_GUARD_BOTH_INT  
_BINARY_OP_ADD_INT
```

```
_SET_IP  
STORE_FAST  
STORE_FAST
```

```
_SET_IP  
_JUMP_TO_TOP
```

Background

CPython 3.13: Micro-Op Traces (`-X uops`)

```
_SET_IP  
_ITER_CHECK_RANGE  
_IS_ITER_EXHAUSTED_RANGE  
_POP_JUMP_IF_TRUE  
_ITER_NEXT_RANGE
```

```
_SET_IP  
STORE_FAST
```

```
_SET_IP  
LOAD_FAST  
LOAD_FAST
```

```
_SET_IP  
LOAD_FAST  
  
_SET_IP  
_GUARD_BOTH_INT  
_BINARY_OP_ADD_INT
```

```
_SET_IP  
STORE_FAST  
STORE_FAST
```

```
_SET_IP  
_JUMP_TO_TOP
```

Background

CPython 3.13: Optimized Micro-Op Traces

```
_SET_IP  
_ITER_CHECK_RANGE  
_IS_ITER_EXHAUSTED_RANGE  
_POP_JUMP_IF_TRUE  
_ITER_NEXT_RANGE
```

```
_SET_IP  
STORE_FAST
```

```
_SET_IP  
LOAD_FAST  
LOAD_FAST
```

```
_SET_IP  
LOAD_FAST
```

```
_SET_IP  
_GUARD_BOTH_INT  
_BINARY_OP_ADD_INT
```

```
_SET_IP  
STORE_FAST  
STORE_FAST
```

```
_SET_IP  
_JUMP_TO_TOP
```

Background

CPython 3.13: Optimized Micro-Op Traces

```
_SET_IP  
_ITER_CHECK_RANGE  
_IS_ITER_EXHAUSTED_RANGE  
_POP_JUMP_IF_TRUE  
_ITER_NEXT_RANGE
```

```
_SET_IP  
STORE_FAST
```

```
_SET_IP  
LOAD_FAST  
LOAD_FAST
```

```
_SET_IP  
LOAD_FAST  
  
_SET_IP  
_GUARD_BOTH_INT  
_BINARY_OP_ADD_INT
```

```
_SET_IP  
STORE_FAST  
STORE_FAST
```

```
_SET_IP  
_JUMP_TO_TOP
```

Background

CPython 3.13: Optimized Micro-Op Traces

```
_SET_IP  
_ITER_CHECK_RANGE  
_IS_ITER_EXHAUSTED_RANGE  
_POP_JUMP_IF_TRUE  
_ITER_NEXT_RANGE
```

```
STORE_FAST
```

```
LOAD_FAST  
LOAD_FAST
```

```
LOAD_FAST
```

```
_SET_IP  
_GUARD_BOTH_INT  
_BINARY_OP_ADD_INT
```

```
STORE_FAST  
STORE_FAST
```

```
_SET_IP  
_JUMP_TO_TOP
```

Background

CPython 3.13: Optimized Micro-Op Traces

`_SET_IP`
`_ITER_CHECK_RANGE`
`_IS_ITER_EXHAUSTED_RANGE`
`_POP_JUMP_IF_TRUE`
`_ITER_NEXT_RANGE`

`STORE_FAST`

`LOAD_FAST`
`LOAD_FAST`

`LOAD_FAST`

`_SET_IP`
`_GUARD_BOTH_INT`
`_BINARY_OP_ADD_INT`

`STORE_FAST`
`STORE_FAST`

`_SET_IP`
`_JUMP_TO_TOP`

Background

CPython 3.13: Optimized Micro-Op Traces

	LOAD_FAST
_IS_ITER_EXHAUSTED_RANGE	
_POP_JUMP_IF_TRUE	
_ITER_NEXT_RANGE	
	_BINARY_OP_ADD_INT
STORE_FAST	
	STORE_FAST
	STORE_FAST
LOAD_FAST	
LOAD_FAST	
	_SET_IP
	_JUMP_TO_TOP

Background

CPython 3.13: Optimized Micro-Op Traces

`_IS_ITER_EXHAUSTED_RANGE`

`_POP_JUMP_IF_TRUE`

`_ITER_NEXT_RANGE`

`STORE_FAST`

`LOAD_FAST`

`LOAD_FAST`

`LOAD_FAST`

`_BINARY_OP_ADD_INT`

`STORE_FAST`

`STORE_FAST`

`_SET_IP`

`_JUMP_TO_TOP`

Just-In-Time Compilation

Just-In-Time Compilation

- Technical goals:
 - Remove interpretive overhead
 - Statically compile optimized traces
 - Reduce indirection:
 - "Burn in" constants, caches, and arguments
 - Move data off of frames and into registers
 - Bring hot code paths in-line
- Deployment goals:
 - Broad platform support
 - Few runtime dependencies
 - Low implementation complexity

Just-In-Time Compilation

- Technical goals:
 - Remove interpretive overhead
 - Statically compile optimized traces
 - Reduce indirection:
 - "Burn in" constants, caches, and arguments
 - Move data off of frames and into registers
 - Bring hot code paths in-line
- Deployment goals:
 - Broad platform support
 - Few runtime dependencies
 - Low implementation complexity

Just-In-Time Compilation

Copy-And-Patch Compilation

Copy-And-Patch Compilation

- Haoran Xu and Fredrik Kjolstad. 2021. Copy-and-Patch Compilation: A Fast Compilation Algorithm for High- Level Languages and Bytecode. Proc. ACM Program. Lang. 5, OOPSLA, Article 136 (October 2021), 30 pages. <https://doi.org/10.1145/3485513>
- Haoran Xu. 2023. Building a baseline JIT for Lua automatically. (12 March 2023). Retrieved from <https://sillycross.github.io/2023/05/12/2023-05-12/>.
- A way of automatically turning a C interpreter into a fast template JIT compiler

Copy-And-Patch Compilation

- Compared to WebAssembly baseline compiler (`Liftoff`):
 - 5x faster code generation
 - 50% faster code
- Compared to traditional JIT toolchain (`LLVM -O0`):
 - 100x faster code generation
 - 15% faster code
- Compared to an optimizing JIT with hand-written assembly (`LuaJIT`):
 - Faster on 13/44 benchmarks
 - Only 35% slower overall

Copy-And-Patch Compilation

- At runtime, walk over a sequence of bytecode instructions.
- For each:
 - Copy some static, pre-compiled machine code into executable memory
 - Patch up instructions that need to have runtime data encoded into them

Copy-And-Patch Compilation

- At runtime, walk over a sequence of bytecode instructions.
- For each:
 - Copy some static, pre-compiled machine code into executable memory
 - Patch up instructions that need to have runtime data encoded into them

Copy-And-Patch Compilation

- Copy some static, pre-compiled machine code into executable memory
- Patch up instructions that need to have runtime data encoded into them

Copy-And-Patch Compilation

- When linking or loading a relocatable object file (ELF, COFF, Mach-O, etc.):
 - Copy some static, pre-compiled machine code into executable memory
 - Patch up instructions that need to have runtime data encoded into them

Copy-And-Patch Compilation

```
case LOAD_FAST:
    PyObject *value = frame->localsplus[oparg];
    Py_INCREF(value);
    *stack_pointer++ = value;
    break;
```

Copy-And-Patch Compilation

```
PyObject *value = frame->localsplus[oparg];  
Py_INCREF(value);  
*stack_pointer++ = value;
```

Copy-And-Patch Compilation

```
int MAGICALLY_INSERT_THE_OPARG;
int MAGICALLY_CONTINUE_EXECUTION(_PyInterpreterFrame *frame,
                                   PyObject **stack_pointer);

int
load_fast(_PyInterpreterFrame *frame, PyObject **stack_pointer)
{
    int oparg = MAGICALLY_INSERT_THE_OPARG;
    PyObject *value = frame->localsplus[oparg];
    Py_INCREF(value);
    *stack_pointer++ = value;
    return MAGICALLY_CONTINUE_EXECUTION(frame, stack_pointer);
}
```

Copy-And-Patch Compilation

```
int MAGICALLY_INSERT_THE_OPARG;
int MAGICALLY_CONTINUE_EXECUTION(_PyInterpreterFrame *frame,
                                  PyObject **stack_pointer);

int
load_fast(_PyInterpreterFrame *frame, PyObject **stack_pointer)
{
    int oparg = MAGICALLY_INSERT_THE_OPARG;
    PyObject *value = frame->localsplus[oparg];
    Py_INCREF(value);
    *stack_pointer++ = value;
    return MAGICALLY_CONTINUE_EXECUTION(frame, stack_pointer);
}
```

Copy-And-Patch Compilation

```
extern int MAGICALLY_INSERT_THE_OPARG;
extern int MAGICALLY_CONTINUE_EXECUTION(_PyInterpreterFrame *frame,
                                         PyObject **stack_pointer);

int
load_fast(_PyInterpreterFrame *frame, PyObject **stack_pointer)
{
    int oparg = &MAGICALLY_INSERT_THE_OPARG;
    PyObject *value = frame->localsplus[oparg];
    Py_INCREF(value);
    *stack_pointer++ = value;
    __attribute__((musttail))
    return MAGICALLY_CONTINUE_EXECUTION(frame, stack_pointer);
}
```


Copy-And-Patch Compilation

```
00: 48 b8 00 00 00 00 00 00 00 00 00 movabsq $0x0, %rax
0a: 48 98                                cltq
0c: 49 8b 44 c5 48                      movq     0x48(%r13,%rax,8), %rax
11: 8b 08                                movl     (%rax), %ecx
13: ff c1                                incl     %ecx
15: 74 02                                je       0x19 <load_fast+0x19>
17: 89 08                                movl     %ecx, (%rax)
19: 48 89 45 00                          movq     %rax, (%rbp)
1d: 48 83 c5 08                          addq     $0x8, %rbp
21: e9 00 00 00 00                      jmp      0x26 <load_fast+0x26>
```

```
02: R_X86_64_64      MAGICALLY_INSERT_THE_OPARG
```

```
22: R_X86_64_PLT32   MAGICALLY_CONTINUE_EXECUTION - 0x4
```

Copy-And-Patch Compilation

```
00: 48 b8 00 00 00 00 00 00 00 00 00 movabsq $0x0, %rax
0a: 48 98                                cltq
0c: 49 8b 44 c5 48                      movq     0x48(%r13,%rax,8), %rax
11: 8b 08                                movl     (%rax), %ecx
13: ff c1                                incl     %ecx
15: 74 02                                je       0x19 <load_fast+0x19>
17: 89 08                                movl     %ecx, (%rax)
19: 48 89 45 00                          movq     %rax, (%rbp)
1d: 48 83 c5 08                          addq     $0x8, %rbp
21: e9 00 00 00 00                      jmp      0x26 <load_fast+0x26>
```

```
02: R_X86_64_64      MAGICALLY_INSERT_THE_OPARG
```

```
22: R_X86_64_PLT32   MAGICALLY_CONTINUE_EXECUTION - 0x4
```

Copy-And-Patch Compilation

```
00: 48 b8 00 00 00 00 00 00 00 00 movabsq $0x0, %rax
0a: 48 98                               cltq
0c: 49 8b 44 c5 48                     movq    0x48(%r13,%rax,8), %rax
11: 8b 08                               movl    (%rax), %ecx
13: ff c1                               incl    %ecx
15: 74 02                               je      0x19 <load_fast+0x19>
17: 89 08                               movl    %ecx, (%rax)
19: 48 89 45 00                         movq    %rax, (%rbp)
1d: 48 83 c5 08                         addq    $0x8, %rbp
21: e9 00 00 00 00                     jmp     0x26 <load_fast+0x26>
```

```
02: R_X86_64_64      MAGICALLY_INSERT_THE_OPARG
```

```
22: R_X86_64_PLT32   MAGICALLY_CONTINUE_EXECUTION - 0x4
```

Copy-And-Patch Compilation

```
static const unsigned char LOAD_FAST_bytes[38] = {
    48,    b8,    00,    00,    00,    00,    00,    00,
    00,    00,    48,    98,    49,    8b,    44,    c5,
    48,    8b,    08,    ff,    c1,    74,    02,    89,
    08,    48,    89,    45,    00,    48,    83,    c5,
    08,    e9,    00,    00,    00,    00,
};
static const Hole LOAD_FAST_holes[2] = {
    { 02, R_X86_64_64,    MAGICALLY_INSERT_THE_OPARG,    +0x0},
    { 22, R_X86_64_PLT32, MAGICALLY_CONTINUE_EXECUTION, -0x4},
};
```

Copy-And-Patch Compilation

```
static const unsigned char LOAD_FAST_bytes[38] = {
    0x48, 0xb8, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00,
    0x00, 0x00, 0x48, 0x98, 0x49, 0x8b, 0x44, 0xc5,
    0x48, 0x8b, 0x08, 0xff, 0xc1, 0x74, 0x02, 0x89,
    0x08, 0x48, 0x89, 0x45, 0x00, 0x48, 0x83, 0xc5,
    0x08, 0xe9, 0x00, 0x00, 0x00, 0x00,
};

static const Hole LOAD_FAST_holes[2] = {
    {0x02, R_X86_64_64,      MAGICALLY_INSERT_THE_OPARG,    +0x0},
    {0x22, R_X86_64_PLT32,  MAGICALLY_CONTINUE_EXECUTION, -0x4},
};
```

brandtbucher/cpython

brandtbucher/cpython
justin

`github.com/brandtbucher/cpython/tree/justin`

Copy-And-Patch Compilation

- Build time:
 - ~700 lines of complex Python
 - ~100 lines of complex C
 - LLVM dependency
- Run time:
 - ~300 lines of simple C (hand-written)
 - ~3000 lines of simple C (generated)
 - No dependencies

Copy-And-Patch Compilation

- Build time:
 - ~700 lines of complex Python
 - ~100 lines of complex C
 - LLVM dependency
- Run time:
 - ~300 lines of simple C (hand-written)
 - ~3000 lines of simple C (generated)
 - No dependencies

Platform Support

Platform Support

Microsoft Windows

- `i686-pc-windows-msvc/msvc`
- `x86_64-pc-windows-msvc/msvc`

Platform Support

All Tier One Platforms

- `i686-pc-windows-msvc/msvc`
- `x86_64-apple-darwin/clang`
- `x86_64-pc-windows-msvc/msvc`
- `x86_64-unknown-linux-gnu/gcc`

Platform Support

All Tier One & Tier Two Platforms

- `aarch64-apple-darwin/clang`
- `aarch64-unknown-linux-gnu/clang`
- `aarch64-unknown-linux-gnu/gcc`
- `i686-pc-windows-msvc/msvc`
- `x86_64-apple-darwin/clang`
- `x86_64-pc-windows-msvc/msvc`
- `x86_64-unknown-linux-gnu/clang`
- `x86_64-unknown-linux-gnu/gcc`

Platform Support

All Tier One & Tier Two Platforms (Except PowerPC)

- `aarch64-apple-darwin/clang`
- `aarch64-unknown-linux-gnu/clang`
- `aarch64-unknown-linux-gnu/gcc`
- `i686-pc-windows-msvc/msvc`
- `powerpc64le-unknown-linux-gnu/gcc`
- `x86_64-apple-darwin/clang`
- `x86_64-pc-windows-msvc/msvc`
- `x86_64-unknown-linux-gnu/clang`
- `x86_64-unknown-linux-gnu/gcc`

Platform Support

All Tier One & Tier Two Platforms

- `aarch64-apple-darwin/clang`
- `aarch64-unknown-linux-gnu/clang`
- `aarch64-unknown-linux-gnu/gcc`
- `i686-pc-windows-msvc/msvc`
- `x86_64-apple-darwin/clang`
- `x86_64-pc-windows-msvc/msvc`
- `x86_64-unknown-linux-gnu/clang`
- `x86_64-unknown-linux-gnu/gcc`

Platform Support

All Tier One & Tier Two Platforms (Including Cross-Compiles!)

- `aarch64-apple-darwin/clang`
- `aarch64-unknown-linux-gnu/clang`
- `aarch64-unknown-linux-gnu/gcc`
- `i686-pc-windows-msvc/msvc`
- `x86_64-apple-darwin/clang`
- `x86_64-pc-windows-msvc/msvc`
- `x86_64-unknown-linux-gnu/clang`
- `x86_64-unknown-linux-gnu/gcc`

